

Lock-free Binary Search Tree and Skip List

Chenhao Li

April 29, 2022



① Background

② Implementation

③ Results

④ Conclusion

1 Background

2 Implementation

3 Results

4 Conclusion

Dictionary

- Dictionary is an important abstract data type in real-world applications.
 - search, insert, remove.
- Binary search tree and skip list are efficient implementations for the dictionary when a partial order relationship is defined.
- It is important (and interesting) and to build lock-free binary search tree and skip list.

Binary Search Tree and Skip List

- BST associates each key with a node and organizes them in a tree structure. The key in the parent node should be strictly greater than keys in the left subtree, and strictly smaller than keys in the right subtree.
- Skip list consists of multiple layers of ordered linked lists. Each node stores a set of `next` pointers, whose number is determined by an exponential distribution. The bottom layer contains all nodes, and higher layers skip multiple nodes.

① Background

② Implementation

③ Results

④ Conclusion

Lock-free Binary Search Tree

- It does not strictly follow the original BST definition.
- Leaf nodes store all keys, internal nodes store parts of keys for routing. Leaf nodes have zero children, internal nodes have two children.
- Use the least significant two bits of child pointers for marking.
 - flag: both the parent and this child are going to be removed.
 - tag: the parent is going to be removed from the tree.
- Once an edge has been flagged or tagged, it cannot point to another node.
- Please refer to [2] for more details.

Lock-free Skip List

- Use the least significant bit in each level of next pointers for marking.
- The mark bit is set before physically removing the node to prevent concurrent operations from inserting after this node and losing the inserted node.
- Search physically remove marked nodes between a level's predecessor and successor.
- Please refer to [1] for more details.

Garbage Collection

- Garbage Collection is non-trivial for lock-free data structures in a language without builtin GC.
- When a thread decides a node can be freed, it may still be reachable in another thread.
- ABA problem: a pointer is freed and reclaimed by the memory allocator, causing CAS instructions to make a wrong judgment that pointer fields have not changed.

Garbage Collection

- Use an atomic integer to represent the current state.
 - The least significant bit: whether there is a pending GC.
 - Higher bits count the number of active operations.
- Push freed pointers into a private list to the thread. When any operation ends, if the number of freed pointers reaches a threshold, set the pending bit with CAS.
- When any operation starts, wait until pending is cleared.
- The thread successfully setting pending waits until there is no active operations and performs GC by freeing all pointers in all threads, and clear pending.

① Background


② Implementation

③ Results

④ Conclusion

Performance

- 1000000 random operations. Modification (insert+remove) ratios are 0.01, 0.1, 0.5, 0.9 respectively.




Performance

- LockedSet is the fastest when parallelism is low and modifications are very few.
- Tree and SkipList can benefit from multithreading well, but more modifications also limit their acceleration.
- Overall, Tree is slightly faster than SkipList (average 15.88%).


Memory Usage

- Intercept malloc/free and overload new/delete.
- Tree > LockedSet > SkipList in terms of memory usage.
- Some jitter in Tree and SkipList due to GC.



Imbalance Insertion

- Insert partially sorted data into dictionaries. Measure sortedness by the longest distance D between reverse pairs.
- LockedSet and SkipList are resistant to sorted data.
- Tree is significantly affected by the degree of sortedness.
- Parallel inserts may break the ordering of input data.



1 Background

2 Implementation

3 Results

4 Conclusion

- If parallelism is low and the main workload is read, A simple STL data structure protected by a readers-writer lock gives the best performance.
- If parallelism is heavy, the lock-free binary search tree has slightly better performance while introducing many limitations.
- The lock-free skip list trades flexibilities for slightly lower performance, and is more useful in most systems.
- The amount of code of the two lock-free data structures is roughly the same, and the skip list is slightly less.

- [1] K. Fraser. Practical lock-freedom. Technical report, University of Cambridge, Computer Laboratory, 2004.
- [2] A. Natarajan and N. Mittal. Fast concurrent lock-free binary search trees. In *Proceedings of the 19th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, PPOPP '14, page 317328, New York, NY, USA, 2014. Association for Computing Machinery. ISBN 9781450326568. doi: 10.1145/2555243.2555256. URL <https://doi.org/10.1145/2555243.2555256>.

Thanks!